

DATA PATH IN COMPUTER SYSTEMS

YUZHE TANG

CF: OSPP-11.4

PAGE FAULT HANDLING (3.4)

DISK READ

CALLER OF DISK READ

- `syscall read`
 0. os pauses caller process of `read ()` to a wait queue
 1. os calls the code procedure of **disk read**
 2. os's interrupt handler then copies the data from kernel's buffer into the process's (virtual) address space.
 3. remove `p1` from wait queue; when `p1` gets scheduled next time, it's returned from `read ()`
- page-fault handler
 0. os pauses caller process of `read ()` to a wait queue
 1. os calls the code procedure of **disk read**
 2. os update page-table and return page-fault handler.

DISK READ (OS PERSPECTIVE)

- d1. mmio: os uses memory-mapped IO to:
 - tell the disk to read the requested data
 - to set up DMA
- d2. DMA: disk then transfer the data to its internal memory before DMAing it to main memory (kernel's buffer).
- d3. notify DMA completion.

1. MMIO MEMORY-MAPPED IO

1. IO bus:

- one end connected to control registers in device (disk)
 - device control registers used to transmit data/commands to device
- another end connected to system's memory bus

2. mmio: map device reg to memory addr space, allowing CPU to access it like regular memory blocks

- e.g. `MOV rax, 0xC0003000`: move value of `rax` to disk control reg.
- versus port-mapped io: pmio uses dedicated instructions `IN/OUT` (x86), now lessly used.

- physical address space is interpreted by hardware; diff. regions assigned to diff. devices.
 - in 32-bit addr. space, a little over 2GB are regions assigned to IO devices.

	VAddr Space	Page table	PAddr Space	
	+-----+ kernel	+-----+ 	+-----+ DRAM	0X00000000
fixed size	+-----+ .text	+-----+ 	+-----+ 	
	+-----+ data	+-----+ 	+-----+ 	0X7FFFFFFF
	+-----+ heap	+-----+ 	+-----+ AUDIO CTRL	0XC0001000
dyn size	+-----+ 	+-----+ 	+-----+ KEYBOARD C	0XC0002000
	+-----+ stack	+-----+ 	+-----+ DISK CTRL	0XC0003000
	+-----+		+-----+	

D2.DMA

- disk io is suuuuuper slow, can't afford to go to CPU for every word transfer
 - instead, disk io transfer words in batch, and w.o. CPU to intervene
- DMA transfers data block from disk's internal memory to system's main memory (target page)
- DMA requires os to "pin" target pages in physical memory; can't be swapped out (e.g. by paging)

D3. INTERRUPT VS POLLING

- when DMA is done, need a mechanism to notify os:
 - "hey CPU! your requested data is now ready in kernel buffer"
- CPU polling the control reg is slow (causing bus IO!)
- raising an interrupt is faster